Computer Networks
COL 334/672

Congestion Control

Tarun Mangla

Slides adapted from KR

Sem 1, 2024-25

Recap: TCP Congestion Control

- End-to-end congestion control algorithms (CCAs)
- Classic CCAs: TCP Reno, TCP Vegas
  - Additive Increase, Multiplicative Decrease (AIMD)
- Slow in case of "long, fat pipes" or networks with high bandwidth-delay product
- TCP CUBIC
  - Increase fast when further away from cwnd where last loss occurred
  - Increase slowly when around cwnd where last loss occurred

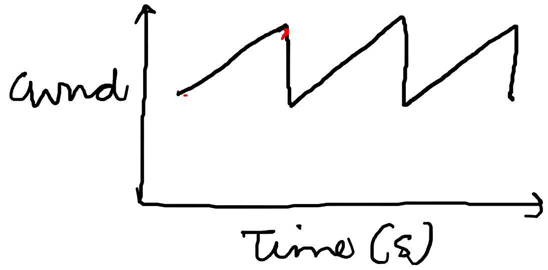

Limitations of a Loss-based CCA

→ Not all losses are due to congestion

- Relying on loss to detect congestion is too reactive

→ losses occure after buffers are full [Bufferbloat] →

Why?

- Waits for queues to build up in the router

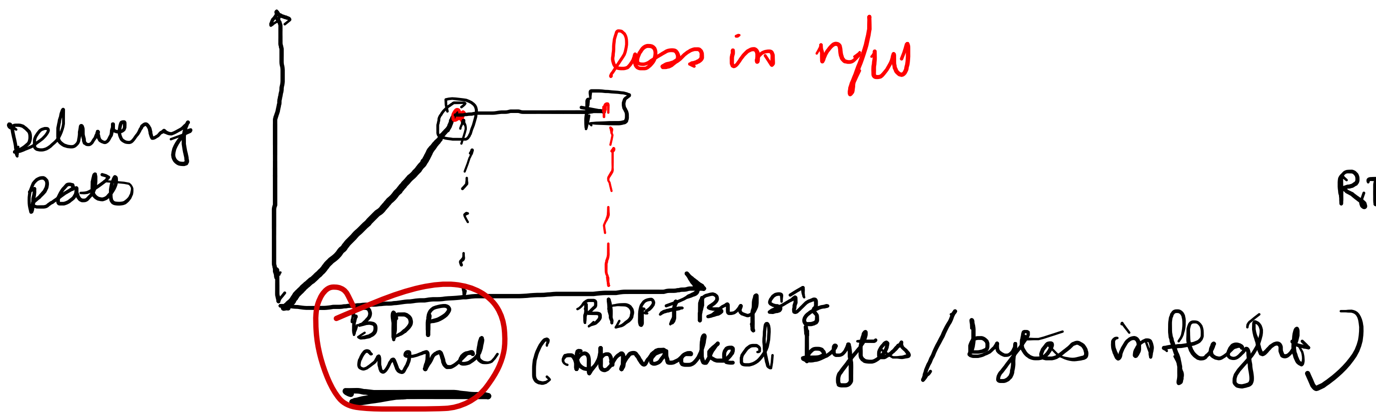

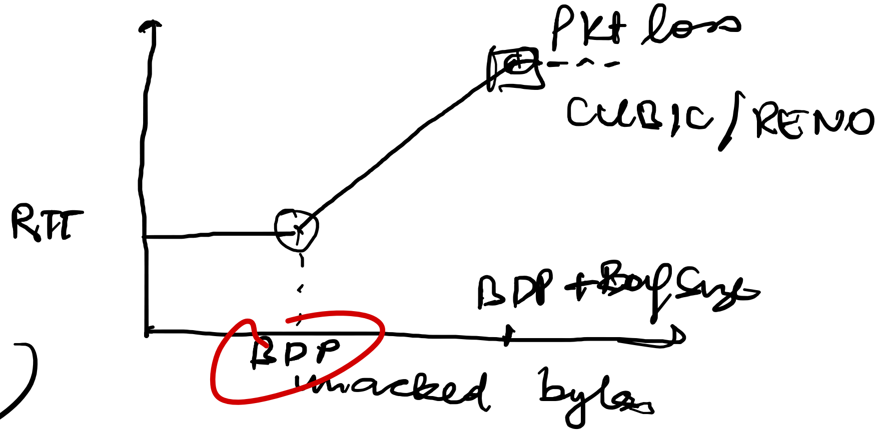

Can we think of another signal for detecting congestion?

Increase in RTTs

Delay-based TCP congestion control

Keeping sender-to-receiver pipe "just full enough, but no fuller": keep bottleneck link busy transmitting, but avoid high delays/buffering

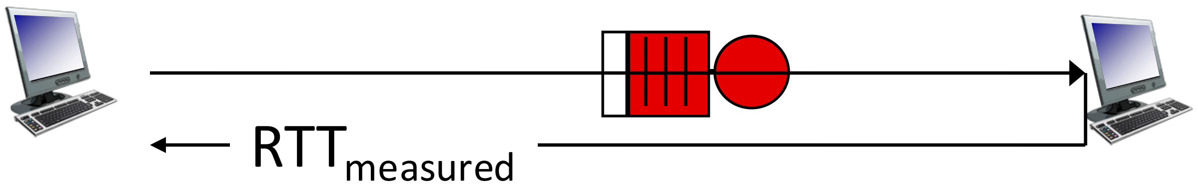

$|T_{measure} - T_{uncongested}| < \alpha$

Action: $\uparrow$ cwnd (linearly)

One Example – TCP Vegas

- RTT_{min} - minimum observed RTT (uncongested path)
- measured throughput: $\dfrac{\text{\# bytes sent in last RTT interval}}{RTT_{measured}}$
- uncongested throughput: $\dfrac{\text{\# bytes sent in last RTT interval}}{RTT_{min}}$

$|T_{measure} - T_{uncongested}| > \beta$

Action: $\downarrow$ cwnd (multiplicatively)

$\alpha \le diff < \beta$ Action: same cwnd

(1) RTT calculations can be unreliable

Challenge with Delay-based CCAs

TCP BBR → Google
(Bottleneck B/w
RTT)

- Don't interact well with loss-based CCAs

- What happens when a delay-based CCA competes with a loss-based CCA?

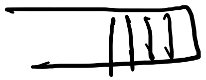

delay based CCA ↓ cwnd

loss-based CCA wait for loss to happen

- Other limitations?

Network-assisted Congestion Control

- Routers in the network help in congestion control
- What are the possible approaches?
  - Tell end points about congestion   Explicit Congestion Notification
  - Manage router buffer but let end-points figure   Active queue management

Explicit congestion notification (ECN)

8 -bits → 2 bits

→ Data center Networks

(DCTCP)

↓

ECN notification

Use header in both Network and Transport Layer

- two bits in IP header (ToS field) marked *by network router* to indicate congestion
- congestion indication carried to destination
- destination sets ECE bit on ACK segment to notify sender of congestion
- sender reduces the congestion window on receiving an ACK with ECE bit set
- Limitation: Requires support from all router in the network path

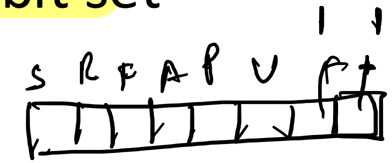

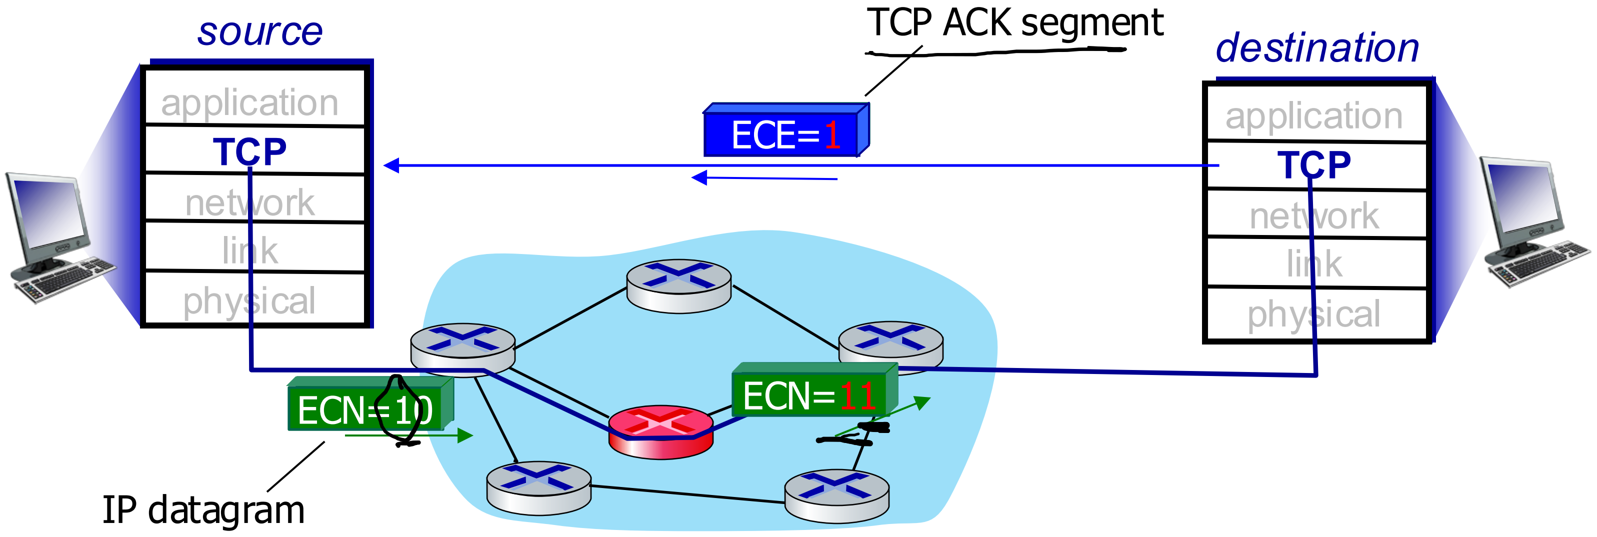

Active Queue Management

- Routers actively control the buffer queues to indirectly aid congestion control
- Why routers? Routers can most accurately identify queuing delays
- Any AQM techniques? Fair queuing, weighted fair queuing

Queuing Discipline

FQ

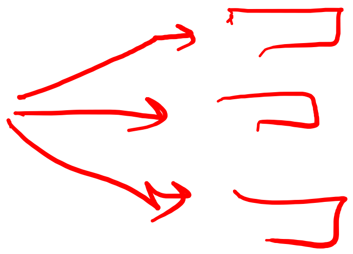

AQM Examples

Drop

- **Random Early Detection (RED)**

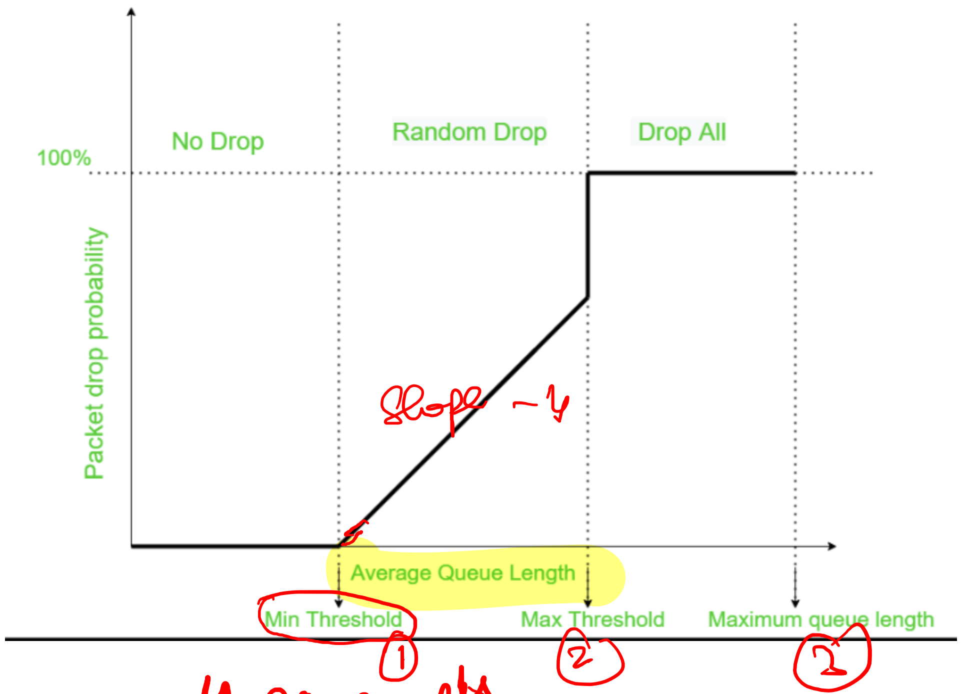

4 parameters

- **CoDel (Controlled Delay)**

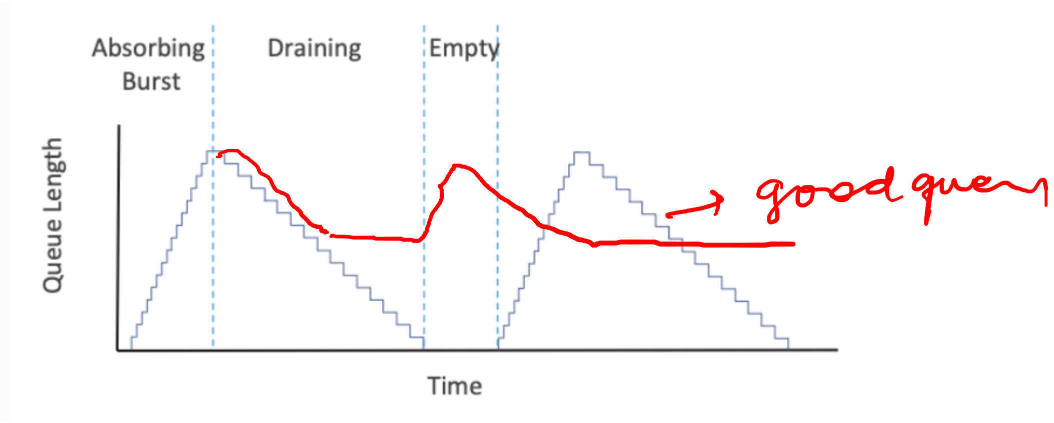

- **Intuition**: CoDel largely ignore queues that last less than an RTT, but starts taking action as soon as a queue persists for more than an RTT

queuing delay > T_1
&
happens for T_2 time

Attendance